

transformer_for_machine_translation

July 15, 2024

```
[ ]: import math

import torch
import torch.nn as nn
from torch.utils.data import dataset
import torch.nn.functional as F

import numpy as np
import matplotlib.pyplot as plt

[ ]: class MultiHeadAttention(nn.Module):
    def __init__(self, d_k, d_model, n_heads, max_len, causal = False):
        super().__init__()

        self.d_k = d_k
        self.n_heads = n_heads

        self.key = nn.Linear(d_model, d_k*n_heads)
        self.query = nn.Linear(d_model, d_k*n_heads)
        self.value = nn.Linear(d_model, d_k*n_heads)

        self.fc = nn.Linear(d_k*n_heads, d_model)

        self.causal = causal
        if causal:
            cm = torch.tril(torch.ones(max_len, max_len))
            self.register_buffer(
                "causal_mask",
                cm.view(1,1,max_len,max_len)
            )

    def forward(self, q, k, v, pad_mask = None):
        q = self.query(q)
        k = self.key(k)
        v = self.value(v)

        N = q.shape[0]
        T_output = q.shape[1]
```

```

T_input = k.shape[1]

q = q.view(N,T_output,self.n_heads, self.d_k).transpose(1,2)
k = k.view(N,T_input,self.n_heads, self.d_k).transpose(1,2)
v = v.view(N,T_input,self.n_heads, self.d_k).transpose(1,2)

attn_scores = q@k.transpose(-2,-1)/math.sqrt(self.d_k)
if pad_mask is not None:
    attn_scores = attn_scores.masked_fill(
        pad_mask[:,None,None, :] == 0, float('-inf'))
    )
if self.causal:
    attn_scores = attn_scores.masked_fill(
        self.causal_mask[:, :, :T_output, :T_input] == 0, float('-inf'))
    )
attn_weights = F.softmax(attn_scores, dim = -1)

A = attn_weights@v

A = A.transpose(1,2)
A = A.contiguous().view(N, T_output, self.d_k * self.n_heads)

return self.fc(A)

```

```

[ ]: class EncoderBlock(nn.Module):
    def __init__(self, d_k,d_model, n_heads, max_len, dropout_prob = 0.1):
        super().__init__()

        self.ln1 = nn.LayerNorm(d_model)
        self.ln2 = nn.LayerNorm(d_model)
        self.mha = MultiHeadAttention(d_k, d_model, n_heads, max_len, causal = False)
        self.ann = nn.Sequential(
            nn.Linear(d_model, d_model *4),
            nn.GELU(),
            nn.Linear(d_model*4, d_model),
            nn.Dropout(dropout_prob)
        )
        self.dropout = nn.Dropout(p=dropout_prob)

    def forward(self, x, mask = None):
        x = self.ln1(x+self.mha(x,x,x,mask))
        x = self.ln2(x+self.ann(x))
        x = self.dropout(x)
        return x

```

```

[ ]: class DecoderBlock(nn.Module):
    def __init__(self, d_k,d_model, n_heads,max_len, dropout_prob = 0.1):

```

```

super().__init__()

self.ln1 = nn.LayerNorm(d_model)
self.ln2 = nn.LayerNorm(d_model)
self.ln3 = nn.LayerNorm(d_model)
self.mha1 = MultiHeadAttention(d_k, d_model, n_heads, max_len, causal = True)
self.mha2 = MultiHeadAttention(d_k, d_model, n_heads, max_len, causal =
→False)
self.ann = nn.Sequential(
    nn.Linear(d_model, d_model *4),
    nn.GELU(),
    nn.Linear(d_model*4, d_model),
    nn.Dropout(dropout_prob)
)
self.dropout = nn.Dropout(p=dropout_prob)

def forward(self, enc_output, dec_input, enc_mask = None, dec_mask = None):
    x = self.ln1(dec_input+self.mha1(dec_input,dec_input,dec_input,dec_mask))
    x = self.ln2(x+self.mha2(x,enc_output,enc_output,enc_mask))
    x = self.ln3(x+self.ann(x))
    x = self.dropout(x)
    return x

```

```

[ ]: class PositionalEncoding(nn.Module):
    def __init__(self, d_model, max_len = 2048, dropout_prob = 0.1):
        super().__init__()

        self.dropout = nn.Dropout(p=dropout_prob)

        position = torch.arange(max_len).unsqueeze(1)
        exp_term = torch.arange(0,d_model, 2)
        div_term = torch.exp(exp_term*(-math.log(10000.0)/d_model))
        pe = torch.zeros(1,max_len, d_model)
        pe[0,:,0::2] = torch.sin(position * div_term)
        pe[0,:,1::2] = torch.cos(position * div_term)
        self.register_buffer('pe',pe)

    def forward(self, x):
        x = x+self.pe[:,x.size(1),:]
        return self.dropout(x)

```

```

[ ]: class Encoder(nn.Module):
    def __init__(self,
        vocab_size,
        max_len,
        d_k,
        d_model,

```

```

        n_heads,
        n_layers,

        dropout_prob):
    super().__init__()
    self.embedding = nn.Embedding(vocab_size, d_model)
    self.pos_encoding = PositionalEncoding(d_model, max_len, dropout_prob)
    transformer_blocks = [
        EncoderBlock(
            d_k,
            d_model,
            n_heads,
            dropout_prob) for _ in range(n_layers)]
    self.transformer_blocks = nn.Sequential(*transformer_blocks)
    self.ln = nn.LayerNorm(d_model)

    def forward(self, x, pad_mask = None):
        x = self.embedding(x)
        x = self.pos_encoding(x)
        for block in self.transformer_blocks:
            x = block(x, pad_mask)

        x = self.ln(x)

    return x

```

```

[ ]: class Decoder(nn.Module):
    def __init__(self,
        vocab_size,
        max_len,
        d_k,
        d_model,
        n_heads,
        n_layers,
        dropout_prob):
        super().__init__()
        self.embedding = nn.Embedding(vocab_size, d_model)
        self.pos_encoding = PositionalEncoding(d_model, max_len, dropout_prob)
        transformer_blocks = [
            DecoderBlock(
                d_k,
                d_model,
                n_heads,
                max_len,
                dropout_prob) for _ in range(n_layers)]

```

```

self.transformer_blocks = nn.Sequential(*transformer_blocks)
self.ln = nn.LayerNorm(d_model)
self.fc = nn.Linear(d_model, vocab_size)

def forward(self, enc_output, dec_input, enc_mask = None, dec_mask = None):
    x = self.embedding(dec_input)
    x = self.pos_encoding(x)
    for block in self.transformer_blocks:
        x = block(enc_output, x, enc_mask, dec_mask)

    x = self.ln(x)
    x = self.fc(x)

    return x

```

```

[ ]: class Transformer(nn.Module):
    def __init__(self, encoder, decoder):
        super().__init__()
        self.encoder = encoder
        self.decoder = decoder

    def forward(self, enc_input, dec_input, enc_mask, dec_mask):
        enc_output = self.encoder(enc_input, enc_mask)
        dec_output = self.decoder(enc_output, dec_input, enc_mask, dec_mask)
        return dec_output

```

```

[ ]: #test_it
encoder = Encoder(
    vocab_size = 20_000,
    max_len = 512,
    d_k = 16,
    d_model = 64,
    n_heads = 4,
    n_layers = 2,
    dropout_prob = 0.1
)
decoder = Decoder(
    vocab_size = 10_000,
    max_len = 512,
    d_k = 16,
    d_model = 64,
    n_heads = 4,
    n_layers = 2,
    dropout_prob = 0.1
)
transformer = Transformer(encoder, decoder)

```

```
[ ]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
encoder.to(device)
decoder.to(device)
```

```
[ ]: xe = np.random.randint(0,20_000, size = (8,512))
xe_t = torch.tensor(xe).to(device)

xd = np.random.randint(0,10_000, size = (8,256))
xd_t = torch.tensor(xd).to(device)

maske = np.ones((8,512))
maske[:,256:] = 0
maske_t = torch.tensor(maske).to(device)

maskd = np.ones((8,256))
maskd[:,128:] = 0
maskd_t = torch.tensor(maskd).to(device)

out = transformer(xe_t,xd_t,maske_t,maskd_t)
out.shape
```

```
[ ]: torch.Size([8, 256, 10000])
```

```
[ ]: !head spa.txt
```

```
[ ]: import pandas as pd
df = pd.read_csv('spa.txt', sep = "\t", header = None)
df.head()
```

```
[ ]: df.shape
```

```
[ ]: (115245, 2)
```

```
[ ]: df = df.iloc[:30_000]
```

```
[ ]: df.columns = ['en','es']
df.to_csv('spa.csv', index = None)
```

```
[ ]: !head spa.csv
```

```
[ ]: !pip install transformers datasets sentencepiece sacremoses
```

```
[ ]: from datasets import load_dataset
raw_dataset = load_dataset('csv', data_files = 'spa.csv')
```

```
[ ]: raw_dataset
```

```
[ ]: split = raw_dataset['train'].train_test_split(test_size = 0.3, seed = 42)
split
```

```
[ ]: from transformers import AutoTokenizer

model_checkpoint = "Helsinki-NLP/opus-mt-en-es"
tokenizer = AutoTokenizer.from_pretrained(model_checkpoint)
```

```
[ ]: en_sentence = split['train'][0]['en']
es_sentence = split['train'][0]['es']

inputs = tokenizer(en_sentence)
targets = tokenizer(text_target = es_sentence)

tokenizer.convert_ids_to_tokens(targets['input_ids'])
```

```
[ ]: es_sentence
```

```
[ ]: max_input_length = 128
max_target_length = 128

def preprocess_function(batch):
    model_inputs = tokenizer(
        batch['en'], max_length = max_input_length, truncation = True
    )

    labels = tokenizer(
        text_target = batch['es'], max_length = max_target_length, truncation = True
    )

    model_inputs["labels"] = labels["input_ids"]
    return model_inputs
```

```
[ ]: tokenized_datasets = split.map(
    preprocess_function,
    batched = True,
    remove_columns = split['train'].column_names
)
```

```
[ ]: tokenized_datasets
```

```
[ ]: from transformers import DataCollatorForSeq2Seq

data_collator = DataCollatorForSeq2Seq(tokenizer)
```

```
[ ]: batch = data_collator([tokenized_datasets["train"][i] for i in range(0,5)])  
batch.keys()
```

```
[ ]: batch['input_ids']
```

```
[ ]: batch['attention_mask']
```

```
[ ]: batch['labels']
```

```
[ ]: tokenizer.all_special_ids
```

```
[ ]: [0, 1, 65000]
```

```
[ ]: tokenizer.all_special_tokens
```

```
[ ]: tokenizer('<pad>')
```

```
[ ]: from torch.utils.data import DataLoader
```

```
train_loader = DataLoader(  
    tokenized_datasets["train"],  
    shuffle = True,  
    batch_size = 32,  
    collate_fn = data_collator  
)
```

```
valid_loader = DataLoader(  
    tokenized_datasets["test"],  
    batch_size = 32,  
    collate_fn = data_collator  
)
```

```
[ ]: for batch in train_loader:  
    for k, v in batch.items():  
        print("k:", k, "v.shape", v.shape)  
    break
```

```
[ ]: tokenizer.vocab_size
```

```
[ ]: tokenizer.decode([60000])
```

```
[ ]: tokenizer.add_special_tokens({"cls_token": "<s>"})
```

```
[ ]: tokenizer("<s>")
```

```
[ ]: tokenizer.vocab_size
```



```
[ ]: encoder = Encoder(vocab_size = tokenizer.vocab_size + 1,
                      max_len = 512,
                      d_k = 16,
                      d_model = 64,
                      n_heads = 4,
                      n_layers = 2,
                      dropout_prob = 0.1)
decoder = Decoder(vocab_size = tokenizer.vocab_size + 1,
                  max_len = 512,
                  d_k = 16,
                  d_model = 64,
                  n_heads = 4,
                  n_layers = 2,
                  dropout_prob = 0.1)

transformer = Transformer(encoder, decoder)
```

```
[ ]: encoder.to(device)
      decoder.to(device)
```

```
[ ]: criterion = nn.CrossEntropyLoss(ignore_index = -100)
      optimizer = torch.optim.Adam(transformer.parameters())
```

```
[ ]: from datetime import datetime
def train(model, criterion, optimizer, train_loader, valid_loader, epochs):
    train_losses = np.zeros(epochs)
    test_losses = np.zeros(epochs)

    for it in range(epochs):
        model.train()
        t0 = datetime.now()
        train_loss = []

        for batch in train_loader:
            batch = {k:v.to(device) for k,v in batch.items()}

            optimizer.zero_grad()

            enc_input = batch['input_ids']
            enc_mask = batch['attention_mask']
            targets = batch['labels']

            dec_input = targets.clone().detach()
            dec_input = torch.roll(dec_input, shifts =1, dims =1)
            dec_input[:,0] = 65_001

            dec_input = dec_input.masked_fill(
```

```

        dec_input == -100, tokenizer.pad_token_id
    )

    dec_mask = torch.ones_like(dec_input)
    dec_mask = dec_mask.masked_fill(dec_input == tokenizer.pad_token_id, 0)

    outputs = model(enc_input, dec_input, enc_mask, dec_mask)
    loss = criterion(outputs.transpose(2,1), targets)

    loss.backward()
    optimizer.step()

    train_loss.append(loss.item())

train_loss = np.mean(train_loss)
model.eval()
test_loss = []

for batch in valid_loader:
    batch = {k:v.to(device) for k, v in batch.items()}

    enc_input = batch['input_ids']
    enc_mask = batch['attention_mask']
    targets = batch['labels']

    dec_input = targets.clone().detach()
    dec_input = torch.roll(dec_input, shifts =1, dims =1)
    dec_input[:,0] = 65_001

    dec_input = dec_input.masked_fill(
        dec_input == -100, tokenizer.pad_token_id
    )

    dec_mask = torch.ones_like(dec_input)
    dec_mask = dec_mask.masked_fill(dec_input == tokenizer.pad_token_id, 0)

    outputs = model(enc_input, dec_input, enc_mask, dec_mask)
    loss = criterion(outputs.transpose(2,1), targets)

    test_loss.append(loss.item())

test_loss = np.mean(test_loss)

train_losses[it] = train_loss
test_losses[it] = test_loss

```

```

    dt = datetime.now() - t0
    print(f'Epoch {it+1}/{epochs}, Train Loss: {train_loss:.4f}, Test Loss:␣
→{test_loss: .4f}, Duration:{dt}')
    return train_losses, test_losses

```

```

[ ]: train_losses, test_losses = train(
    transformer, criterion, optimizer, train_loader, valid_loader, epochs = 15
)

```

```

[ ]: input_sentence = split['test'][10]['en']
    input_sentence

```

```

[ ]: enc_input = tokenizer(input_sentence, return_tensors = 'pt')
    enc_input

```

```

[ ]: dec_input_str = '<s>'

    dec_input = tokenizer(text_target = dec_input_str, return_tensors = 'pt')
    dec_input

```

```

[ ]: enc_input.to(device)
    dec_input.to(device)
    output = transformer(
        enc_input['input_ids'],
        dec_input['input_ids'][:, :-1],
        enc_input['attention_mask'],
        dec_input['attention_mask'][:, :-1]
    )
    output

```

```

[ ]: outptu.shape

```

```

[ ]: enc_output = encoder(enc_input['input_ids'], enc_input['attention_mask'])
    enc_output.shape

```

```

[ ]: dec_output = decoder(
    enc_output,
    dec_input['input_ids'][:, :-1],
    enc_input['attention_mask'],
    dec_input['attention_mask'][:, :-1]
)

    dec_output.shape

```

```

[ ]: torch.allclose(output, dec_output)

```

```
[ ]: dec_input_ids = dec_input['input_ids'][:, :-1]
dec_attn_mask = dec_input['attention_mask'][:, :-1]

for _ in range(32):
    dec_output = decoder(
        enc_output,
        dec_input_ids,
        enc_input['attention_mask'],
        dec_attn_mask
    )

    prediction_id = (torch.argmax(dec_output[:, -1, :], axis = -1))

    dec_input_ids = torch.hstack((dec_input_ids, prediction_id.view(1,1)))

    dec_attn_mask = torch.ones_like(dec_input_ids)

    if prediction_id == 0:
        break
```

```
[ ]: tokenizer.decode(dec_input_ids[0])
```

```
[ ]: split['test'][10]['es']
```

```
[ ]: def translate(input_sentence):
    enc_input = tokenizer(input_sentence, return_tensors = 'pt').to(device)
    enc_output = encoder(enc_input['input_ids'], enc_input['attention_mask'])

    dec_input_ids = torch.tensor([[65_001]], device = device)
    dec_attn_mask = torch.ones_like(dec_input_ids, device = device)

    for _ in range(32):
        dec_output = decoder(
            enc_output,
            dec_input_ids,
            enc_input['attention_mask'],
            dec_attn_mask
        )

        prediction_id = (torch.argmax(dec_output[:, -1, :], axis = -1))

        dec_input_ids = torch.hstack((dec_input_ids, prediction_id.view(1,1)))

        dec_attn_mask = torch.ones_like(dec_input_ids)

        if prediction_id == 0:
            break
```

```
translation = tokenizer.decode(dec_input_ids[0,1:])  
print(translation)
```

```
[ ]: translate("How are you?")
```